

Quantum Machine Learning (QML) Classifiers

Objective

The objective of this project is to design, implement, and benchmark a Quantum Machine Learning (QML) classifier for a binary fraud detection problem. The key goals of this project are to:

- Reduce the dataset to a quantum-feasible feature space using Principal Component Analysis (PCA)
- Train strong classical machine learning baselines
- Build a Variational Quantum Classifier (VQC) using Qiskit
- Compare the performance of classical models and quantum models
- Study the capabilities and limitations of QML under NISQ-era constraints, such as:
 - Limited number of qubits
 - Shallow quantum circuits
 - Expensive quantum simulation

This project demonstrates a complete hybrid machine learning pipeline, covering data preprocessing, classical modeling, quantum circuit design, training, and performance comparison.

2. OVERVIEW

This project follows a structured week-wise workflow:

Weeks 1-2

- Load and clean the dataset
- Perform exploratory data analysis (EDA)
- Apply scaling and dimensionality reduction using **PCA**
- Train classical models:
 - Logistic Regression
 - Random Forest
- Visualize classical model performance

Weeks 2-3

- Design the Quantum Circuit
- Build a **4-qubit VQC model** using:
 - ZZFeatureMap
 - Real Amplitudes Ansatz
 - COBYLA optimizer
- Train the QML model on a reduced dataset
- Predict using the trained quantum circuit

Weeks 4–5

- Compare classical vs quantum accuracy
- Interpret performance differences
- Document limitations and future scope

This end-to-end flow reflects a complete hybrid machine learning pipeline.

3. DATA ANALYSIS and PREPROCESSING

3.1 Loading the Dataset

The fraud dataset was loaded using pandas and inspected using:

- `df.head()`
- `df.info()`
- `df.describe()`

3.2 Data Cleaning

- No critical missing values
- All features were numerical
- The target column **fraud** was confirmed as binary

3.3 Class Distribution

A bar plot was used to visualize class imbalance.

This is important because fraud datasets typically have more non-fraud cases.

3.4 Train–Test Split

- Performed using **stratified split**
- Maintains the same fraud ratio in training and testing

3.5 Feature Scaling

Standardization using:

- **StandardScaler**
Ensures PCA and classical models work correctly.

3.6 Dimensionality Reduction (PCA)

Quantum models are currently limited by the number of available qubits. To address this:

- Principal Component Analysis (PCA) was applied.
- The most informative components were retained.

- Reduced-dimensional representations were used for the quantum model while full features were used for classical baselines.
- This step enabled efficient quantum encoding while minimizing information loss.

3.7 Handling Class Imbalance using SMOTE

Fraud detection datasets are inherently highly imbalanced, with fraudulent transactions forming a very small fraction of the total data. Training models directly on such data can lead to biased predictions toward the majority (non-fraud) class.

To address this issue, Synthetic Minority Over-sampling Technique (SMOTE) was applied during the classical machine learning training phase.

Why SMOTE was used

- Prevents models from being biased toward the majority class
- Improves detection of minority (fraud) cases
- Enhances recall and ROC-AUC performance
- Produces more informative decision boundaries

SMOTE Implementation

- SMOTE was applied only on the training data, after train-test split.
- Synthetic samples of the minority class were generated using nearest-neighbor interpolation.
- The test set was left untouched to ensure fair evaluation.

Impact on Model Performance

- Classical models trained with SMOTE showed improved sensitivity to fraud cases.
- Reduction in false negatives was observed.
- Overall ROC-AUC scores improved compared to training without resampling.

SMOTE played a critical role in strengthening the classical baseline, ensuring that comparisons with the quantum model were meaningful and fair.

4. QUANTUM CIRCUIT DESIGN

The quantum model was built using **Qiskit Machine Learning** components.

4.1 Number of Qubits

- Determined by PCA output
- **4 qubits = 4 PCA features**

4.2 Feature Map

ZZFeatureMap

- Encodes classical data into quantum states
- Introduces entanglement between qubits
- Captures non-linear feature interactions

4.3 Ansatz (Variational Circuit)

Real Amplitudes Ansatz

- Rotation + entanglement layers, 2 repetitions (reps=2)
- Hardware-efficient and stable for training

4.4 Backend & Measurements

- Backend: **StatevectorSampler**
- Ensures faster simulation than real hardware
- Computes measurement probabilities required for classification

The circuit is optimized to be shallow and feasible for NISQ-era devices.

5. TRAINING PIPELINE

5.1 Classical Training

Models implemented:

- Logistic Regression
- Random Forest

Both trained on PCA-transformed features.

Accuracy was calculated on test data, then visualized in a bar plot.

5.2 Quantum Training

The QML training loop includes:

1. **Data encoding** through ZZFeatureMap
2. **Parameterized quantum circuit execution**

3. **Classical optimization** of circuit parameters (COBYLA)

4. **Measurement + prediction**

Because quantum simulation is expensive:

- A reduced dataset (first 150 PCA samples) was used for training
- A shallow ansatz was selected
- Low optimizer iterations were used

5.3 QML Prediction

Final quantum predictions were compared with test labels to compute accuracy.

5.4 Model Evaluation

Performance evaluation included:

- Prediction on unseen test data, Accuracy computation.
- ROC-AUC score calculation for robust fraud detection assessment.

These metrics ensured fair comparison with classical models.

6. COMPARATIVE ANALYSIS

Classical Models

- **Random Forest:** best accuracy
- **Logistic Regression:** moderate accuracy

Quantum Model

- VQC achieved **55–60% accuracy**, demonstrating learning but limited by:
 - few qubits
 - shallow circuits
 - no noise modeling
 - classical optimizer inefficiencies


Visualization

A performance comparison plot was created:

Results Comparison

...

	Model	Accuracy	
0	Variational Quantum Classifier	0.60940	
1	Logistic Regression	0.83930	
2	Random Forest	0.91885	
3	MLP Classifier	0.91970	



Classical Accuracy vs Quantum Accuracy

This clearly shows the gap between classical and quantum models.

7. CONCLUSION

This project successfully demonstrates an end-to-end **Quantum + Classical Machine Learning pipeline** for fraud detection.

Key Insights

- Classical models outperform QML due to better optimization and full feature access.
- QML performance is limited by:
 - small qubit count
 - restricted circuit depth
 - costly quantum simulation
- Despite constraints, the VQC learned meaningful patterns.
-

Final Takeaway

Quantum Machine Learning is promising but not yet superior to classical ML for real-world, structured datasets.

However, as quantum hardware improves, deeper circuits and error mitigation could lead to better quantum performance.

Future Scope

- Test model on real IBM Quantum hardware
- Introduce quantum noise models
- Explore data re-uploading feature maps
- Use parameter-shift gradients for better optimization